

2.3. Vector classes Library

Exercises are located at Exercises/3_Vc/

Solutions are located at Exercises/3_Vc/****/****_solution.cpp

To compile and run exercise programs use the line given in the head-comments in the code.

The results given here are obtained on Intel E7-4860 CPU with gcc4.7.3.

Vc Introduction

Vector classes (Vc) is a free software library to ease explicit vectorization of C++ code. It has an intuitive interface and provides portability between different compilers and compiler versions as well as portability between different vector instruction sets. Thus an application written with Vc can be compiled for AVX, SSE, XeonPhi (MIC) and others SIMD instructions¹.

Similar to the SIMD header files, it provides all basic arithmetic operations and functions, and much more in addition. All functionality, which the headers provide for **fvec** can be used similarly with the class **float_v**. The most useful functionalities of Vc are masks and random memory access.

The mask functionality allows conditioning calculations. For example, an absolute value of **a** can be calculated as:

```
float_m mask = ( a < 0 );  a(mask) = - a;
```

Here **mask** saves the comparison result for each entry of **a** in a form, which can be represented like [true;true;false;true] and **operator()** applies the assignment only to those entries, where the mask is true.

The random memory access functionality is provided by the **gather** and **scatter** functions. For example:

gather fills a vector **a** from an array **A** taking elements with indexes stored in a vector **I**:

```
a.gather( A, I ); // A[I] -> a
```

scatter makes the opposite - fills an array entries from a vector:

```
a.scatter( A, I ); // a -> A[I]
```

3_Vc/0_Matrix: description

The Matrix exercise requires to parallelize the square root extraction over a set of float variables arranged in a square matrix using the Vc library. The initial code implements scalar and vector parts using the SIMD header, the space for the Vc implementation is blank. Therefore the initial output shows 0 time for vector calculations and infinite speed up factor, that should be currently ignored.

¹ <http://code.compeng.uni-frankfurt.de/projects/vc>

	Part of the source code of Matrix.cpp
26	<code>float a[N][N] __attribute__((aligned(16)));</code>
27	<code>float c[N][N] __attribute__((aligned(16)));</code>
28	<code>float c_simd[N][N] __attribute__((aligned(16)));</code>
29	<code>float c_simdVc[N][N] __attribute__((aligned(16)));</code>
30	
31	<code>template<typename T> // required calculations</code>
32	<code>T f(T x) {</code>
33	<code> return sqrt(x);</code>
34	<code>}</code>
...	
49	<code>int main() {</code>
50	
51	<code> // fill classes by random numbers</code>
52	<code> for(int i = 0; i < N; i++) {</code>
53	<code> for(int j = 0; j < N; j++) {</code>
54	<code> a[i][j] = float(rand())/float(RAND_MAX); // put a random value, from 0 to</code>
55	<code>1</code>
56	<code> }</code>
57	<code> }</code>
58	<code> /// -- CALCULATE --</code>
59	<code> /// SCALAR</code>
60	<code>TStopwatch timerScalar;</code>
61	<code> for(int ii = 0; ii < NIter; ii++)</code>
62	<code> for(int i = 0; i < N; i++) {</code>
63	<code> for(int j = 0; j < N; j++) {</code>
64	<code> c[i][j] = f(a[i][j]);</code>
65	<code> }</code>
66	<code> }</code>
67	<code> timerScalar.Stop();</code>
68	
69	<code> /// SIMD VECTORS</code>
70	<code>TStopwatch timerSIMD;</code>
71	<code> for(int ii = 0; ii < NIter; ii++)</code>
72	<code> for(int i = 0; i < N; i++) {</code>
73	<code> for(int j = 0; j < N; j+=fvecLen) {</code>
74	<code> fvec &aVec = (reinterpret_cast<fvec*>(a[i][j]));</code>
75	<code> fvec &cVec = (reinterpret_cast<fvec*>(c_simd[i][j]));</code>
76	<code> cVec = f(aVec);</code>
77	<code> }</code>
78	<code> }</code>
79	<code> timerSIMD.Stop();</code>
80	
81	<code> /// Vc</code>
82	<code>TStopwatch timerVc;</code>
83	<code> //TODO write the code using Vc</code>
84	<code> timerVc.Stop();</code>

	Typical output
	Time scalar: 798.745 ms Time headers: 201.086 ms, speed up 3.97215 Time Vc: 0 ms, speed up inf SIMD and scalar results are the same. ERROR! SIMD and scalar results are not the same.

3_Vc/0_Matrix: solution

Since `float_v` is stored in memory in the same way as `fvec` and the function `f(...)` is a template, the part for Vc can be exactly the same, simply the name of the type must be changed:

	Part of the source code of Matrix_solution.cpp
81	<code>/// Vc</code>
82	<code>TStopwatch timerVc;</code>
83	<code>for(int ii = 0; ii < NIter; ii++)</code>
84	<code>for(int i = 0; i < N; i++) {</code>
85	<code>for(int j = 0; j < N; j+=float_v::Size) {</code>
86	<code>float_v &aVec = (reinterpret_cast<float_v&>(a[i][j]));</code>
87	<code>float_v &cVec = (reinterpret_cast<float_v&>(c_simd[i][j]));</code>
88	<code>cVec = f(aVec);</code>
89	<code>}</code>
90	<code>}</code>
91	<code>timerVc.Stop();</code>
	Typical output after solution
	Time scalar: 798.728 ms Time headers: 201.078 ms, speed up 3.97223 Time Vc: 201.079 ms, speed up 3.97221 SIMD and scalar results are the same. SIMD and scalar results are the same.

Since 4 float variables fit into a single SIMD vector, all calculations are done in parallel and no overhead operations are required, the expected speed-up factor should be 4.

3_Vc/1_QuadraticEquation: description

The QuadraticEquation exercise requires to vectorize using Vc the solution of a set of quadratic equations in three different ways, based on three different data setups: (1) Array of Structures (AoS), (2) Structure of Arrays (SoA) (3) Array of Structures of Arrays (AoSoA).

It is recommended to compile the code with `-fno-tree-vectorize` option to prevent auto-vectorization of the scalar code, otherwise the comparison of the vectorized and scalar codes will not be direct.

The input data is given already in the required formats. An elemental structure `DataAOSElement`, which contains parameters of the equations and the resulting maximum root, is declared for the AoS format. `_mm_malloc` function is used to allocate aligned memory.

	Part of the source code of QuadraticEquation.cpp
23	<code>struct DataA0SElement {</code>
24	<code> float a, b, c, // coefficients</code>
25	<code> x; // a root</code>
26	<code>};</code>
27	
28	<code>struct DataA0S {</code>
29	<code> DataA0S(const int N) {</code>
30	<code> data = (DataA0SElement*) _mm_malloc(sizeof(DataA0SElement)*N,16);</code>
31	<code> }</code>
32	<code> ~DataA0S() {</code>
33	<code> if(data) _mm_free(data);</code>
34	<code> }</code>
35	<code> DataA0SElement *data;</code>
36	<code>};</code>

The SoA format is declared as a whole structure, which contains dynamic arrays.

	Part of the source code of QuadraticEquation.cpp
38	<code>struct DataSOA {</code>
39	
40	<code> DataSOA(const int N) {</code>
41	<code> a = (float*) _mm_malloc(sizeof(float)*N,16);</code>
42	<code> b = (float*) _mm_malloc(sizeof(float)*N,16);</code>
43	<code> c = (float*) _mm_malloc(sizeof(float)*N,16);</code>
44	<code> x = (float*) _mm_malloc(sizeof(float)*N,16);</code>
45	<code> }</code>
46	<code> ~DataSOA()</code>
47	<code> {</code>
48	<code> if(a) _mm_free(a);</code>
49	<code> if(b) _mm_free(b);</code>
50	<code> if(c) _mm_free(c);</code>
51	<code> if(x) _mm_free(x);</code>
52	<code> }</code>
53	
54	<code> float *a, *b, *c, // coefficients</code>
55	<code> *x; // a root</code>
56	<code>};</code>

To define the AoSoA format an elemental structure **DataAOSOAElement** is declared to contain information about **float_v::Size** (4) equations, similarly to **DataSOA**. The elemental structures are packed together in the DataAOSOA structure similarly to **DataAOS**. Memory for all information is allocated in one go in **DataAOSOA** to ensure compact data allocation, afterwards the memory is distributed between elemental structures using **SetMemory** function.

	Part of the source code of QuadraticEquation.cpp
58	<code>struct DataA0S0AElement {</code>
59	
60	<code>void SetMemory(float *m) {</code>
61	<code>a = m;</code>
62	<code>b = m + float_v::Size;</code>
63	<code>c = m + 2*float_v::Size;</code>
64	<code>x = m + 3*float_v::Size;</code>
65	<code>}</code>
66	
67	<code>float *a, *b, *c, // coefficients</code>
68	<code>*x; // a root</code>
69	<code>};</code>
70	
71	<code>struct DataA0S0A {</code>
72	<code>DataA0S0A(const int N) {</code>
73	<code>const int NVectors = N/float_v::Size;</code>
74	
75	<code>data = new DataA0S0AElement[NVectors];</code>
76	<code>memory = (float*) _mm_malloc(sizeof(float)*4*N,16);</code>
77	
78	<code>float *memp = memory;</code>
79	<code>for(int i = 0; i < NVectors; i++) {</code>
80	<code>data[i].SetMemory(memp);</code>
81	<code>memp += float_v::Size*4;</code>
82	<code>}</code>
83	<code>}</code>
84	<code>~DataA0S0A() {</code>
85	<code>_mm_free(memory);</code>
86	<code>delete[] data;</code>
87	<code>}</code>
88	
89	<code>float *memory;</code>
90	<code>DataA0S0AElement *data;</code>
91	<code>};</code>

The given structures is filled by the same random data sample and scalar implementation is given in the exercise code.

	Part of the source code of QuadraticEquation.cpp
150	// fill parameters by random numbers
151	for(int i = 0; i < N; i++) {
152	float a = 0.01 + float(rand())/float(RAND_MAX); // put a random value, from 0.01 to 1.01 (a has not to be equal 0)
153	float b = float(rand())/float(RAND_MAX);
154	float c = -float(rand())/float(RAND_MAX);
155	
156	dataScalar.data[i].a = a;
157	dataScalar.data[i].b = b;
158	dataScalar.data[i].c = c;
159	
160	dataSIMD1.data[i].a = a;
161	dataSIMD1.data[i].b = b;
162	dataSIMD1.data[i].c = c;
163	
164	dataSIMD2.a[i] = a;
165	dataSIMD2.b[i] = b;
166	dataSIMD2.c[i] = c;
167	
168	const int nV = i/float_v::Size;
169	const int iV = i%float_v::Size;
170	dataSIMD3.data[nV].a[iV] = a;
171	dataSIMD3.data[nV].b[iV] = b;
172	dataSIMD3.data[nV].c[iV] = c;
173	}
174	
175	/// -- CALCULATE --
176	
177	// scalar calculations
178	TStopwatch timerScalar;
179	for(int io=0; io<NIterOut; io++)
180	for(int i=0; i<N; i++)
181	{
182	float &a = dataScalar.data[i].a;
183	float &b = dataScalar.data[i].b;
184	float &c = dataScalar.data[i].c;
185	float &x = dataScalar.data[i].x;
186	
187	float det = b*b - 4*a*c;
188	x = (-b+sqrt(det))/(2*a);
189	}
190	timerScalar.Stop();
	Typical output
	Time scalar: 567.055 ms
	Time Vc AOS: 0 ms, speed up inf
	Time Vc SOA: 0 ms, speed up inf
	Time Vc AOSOA: 0 ms, speed up inf
	ERROR! SIMD and scalar results are not the same. SIMD part 1.
	ERROR! SIMD and scalar results are not the same. SIMD part 2.
	ERROR! SIMD and scalar results are not the same. SIMD part 3.

3_Vc/1_QuadraticEquation: solution

(1) To vectorize the AoS format one needs to gather data from different instances of the **DataAOSElement** structure and pack it together (in groups by 4) in **float_v** variables. Since the same data (for example, the **a** variable) is placed in different parts of memory, data coping is necessary here. Once data is grouped, the solution is found using exactly the same code as one given in the scalar part. The data is ungrouped back into the output **x** variable.

	Part of the source code of QuadraticEquation_solution.cpp
194	<code>for(int io=0; io<NIterOut; io++)</code>
195	<code>{</code>
196	<code>for(int i=0; i<NVectors; i++)</code>
197	<code>{</code>
198	<code>// copy input data</code>
199	<code>float_v aV;</code>
200	<code>float_v bV;</code>
201	<code>float_v cV;</code>
202	
203	
204	<code>for(int iV=0; iV<float_v::Size; iV++)</code>
205	<code>{</code>
206	<code> aV[iV] = dataSIMD1.data[i*float_v::Size + iV].a;</code>
207	<code> bV[iV] = dataSIMD1.data[i*float_v::Size + iV].b;</code>
208	<code> cV[iV] = dataSIMD1.data[i*float_v::Size + iV].c;</code>
209	<code>}</code>
210	
211	<code>const float_v det = bV*bV - 4*aV*cV;</code>
212	<code>float_v xV = (-bV+sqrt(det))/(2*aV);</code>
213	
214	<code>// copy output data</code>
215	<code>for(int iE=0; iE<float_v::Size; iE++)</code>
216	<code> dataSIMD1.data[i*float_v::Size+iE].x = xV[iE];</code>
217	<code>}</code>
218	<code>}</code>

(2) The second task is vectorization with the SoA data format. Since in SoA similar data is placed near each other, the **reinterpret_cast** can be used. Once the data is reinterpreted, the solution is found using exactly the same code, as the given one in the scalar part.

	Part of the source code of QuadraticEquation_solution.cpp
223	<code>for(int io=0; io<NIterOut; io++)</code>
224	<code>for(int i=0; i<N; i+=float_v::Size)</code>
225	<code>{</code>
226	<code>float_v& aV = (reinterpret_cast<float_v&>(dataSIMD2.a[i]));</code>
227	<code>float_v& bV = (reinterpret_cast<float_v&>(dataSIMD2.b[i]));</code>
228	<code>float_v& cV = (reinterpret_cast<float_v&>(dataSIMD2.c[i]));</code>
229	<code>float_v& xV = (reinterpret_cast<float_v&>(dataSIMD2.x[i]));</code>
230	
231	<code>const float_v det = bV*bV - 4*aV*cV;</code>
232	<code>xV = (-bV+sqrt(det))/(2*aV);</code>
233	<code>}</code>

(3) The third task is vectorization with the AoSoA data format. The reinterpret cast can be used here in the same way as with the SoA format, just the dereference operator must be applied in addition. The calculations part remains the same as well.

	Part of the source code of QuadraticEquation_solution.cpp
238	<code>for(int io=0; io<NIterOut; io++)</code>
239	<code>for(int i=0; i<NVectors; i++)</code>
240	<code>{</code>
241	<code>float_v& aV = *(reinterpret_cast<float_v*>(dataSIMD3.data[i].a));</code>
242	<code>float_v& bV = *(reinterpret_cast<float_v*>(dataSIMD3.data[i].b));</code>
243	<code>float_v& cV = *(reinterpret_cast<float_v*>(dataSIMD3.data[i].c));</code>
244	<code>float_v& xV = *(reinterpret_cast<float_v*>(dataSIMD3.data[i].x));</code>
245	
246	<code>const float_v det = bV*bV - 4*aV*cV;</code>
247	<code>xV = (-bV+sqrt(det))/(2*aV);</code>
248	<code>}</code>

	Typical output after solution
	Time scalar: 566.821 ms
	Time Vc AOS: 217.513 ms, speed up 2.60592
	Time Vc SOA: 142.603 ms, speed up 3.97482
	Time Vc AOSOA: 143.57 ms, speed up 3.94805
	SIMD 1 and scalar results are the same.
	SIMD 2 and scalar results are the same.
	SIMD 3 and scalar results are the same.

The speed up factor of 4 is expected due to vectorization. It is achieved with the SoA and AoSoA data formats. With AoS additional data regrouping is required, that results in the smaller speed up of 2.6.

3_Vc/2_Newton: description

The Newton exercise requires to vectorize the Newton method for numerical solution of a group of equations.

The method can be explained graphically as in the Fig 1. The task is to find intersection of the curve and the X-axis. The algorithm starts with an initial approximation for root x_0 and takes a function value at




Fig. 1. Explanation of the Newton task.

this point. A tangent line is drawn at this point, an intersection of the tangent line with the X-axis is the next approximation x_1 :

$$x_k \equiv x_{k-1} - f(x_{k-1})/f'(x_{k-1})$$

Then the procedure is repeated until the approximation does not change within the required precision.

$$x_k - x_{k-1} < \text{epsilon}$$

A scalar version of the algorithm is given in the exercise. A solution is proposed to perform in two steps. First, vectorize the algorithm, which uses a fixed number of iterations (1000). Then make the number of iterations dependent on the required precision. The vectorized version must give exactly the same result as the scalar one.

	Part of the source code of Newton.cpp
44 45 46 47 48 49 50 51 52	<pre>float FindRootScalar(const float& p1, const float& p2) { float x = 1, x_new = 0; for(; abs((x_new - x)/x_new) > P;) { // for(int i = 0; i < 1000; ++i){ x = x_new; x_new = x - F(x,p1,p2) / Fd(x,p1,p2); } return x_new; }</pre>
	Typical output
	<p>Scalar part: Results are correct! Time: 15.415 ms SIMD part: Results are NOT the same! Time: 0.226021 ms Speed up: 68.2015</p>

3_Vc/2_Newton: solution

The parallelization is achieved by grouping 4 equations together. A complication in this case is that different equations will reach the required precision at different number of iterations. This problem can be solved using Vc masks: the loop exit condition should be stored as a mask. Then the loop is continued until the mask has at least one entry with the value **true**. A mask value should be also used during update of the root approximation in order to reproduce the scalar result.

	Part of the source code of Newton_solution.cpp
56 57 58 59 60 61 62 63 64 65 66 67	<pre> float_v FindRootVector(const float_v& p1, const float_v& p2) { float_v x = 1.f, x_new = 0.f; float_m mask(true); for(; !mask.isEmpty();) { // for(int i = 0; i < 1000; ++i){ x = x_new; x_new(mask) = x - F(x,p1,p2) / Fd(x,p1,p2); mask = abs((x_new - x)/x_new) > P; } return x_new; } </pre>
	Typical output after solution
	<pre> Scalar part: Results are correct! Time: 15.3401 ms SIMD part: Results are the same! Time: 4.49395 ms Speed up: 3.4135 </pre>

A typical speed up factor of vectorization is 4, but since mask is used, the parallelization is not full and at the final iterations only a part of the SIMD vector entries is used. Therefore, the speed up factor is about 3.5.

3_Vc/3_RandomAccess: description

The RandomAccess exercise requires to use different forms of the **gather** and **scatter** functions. For that an input array of **float** is provided and randomly filled. Also an array of random indices is provided. It is required to (1) gather data from the **input** array to the tmp **float_v** variable according to the **index** array; (2) similarly, gather the data, but only when it satisfies a given condition; (3) put the data from the tmp **float_v** variable to the **output** array, when it satisfies a given condition.

Part of the source code of RandomAccess.cpp

```

21 float input[N];
22 float output[N];
23
24 int main() {
25
26     unsigned int index[float_v::Size];
27
28     // fill input array with random numbers from 0 to 1
29     for( int i = 0; i < N; i++ ) {
30         input[i] = float(rand())/float(RAND_MAX);
31     }
32
33     // fill output array with 0
34     for( int i = 0; i < N; i++ ) {
35         output[i] = 0;
36     }
37
38     // fill indices with random numbers from 0 to N-1
39     for( int i = 0; i < float_v::Size; i++ ) {
40         index[i] = static_cast<unsigned int>(float(rand())/float(RAND_MAX)*N);
41     }
42
43     cout << "Indices: ";
44     for(int i=0; i<float_v::Size; i++)
45         cout << index[i] << " ";
46     cout << endl;
47
48     /// gather without masking
49     float_v tmp;
50     //TODO0 gather data with indices "index" from the array "input" into float_v
51     tmp
52     // Use void gather (const float *array, const uint_v &indexes)
53
54     ...
55
56     /// gather with masking
57     float_v tmp2;
58     //TODO0 gather data with indices "index" from the array "input" into float_v
59     tmp2, if the value of "input" large then 0.5
60     // Use void gather (const float *array, const uint_v &indexes, const float_m
61     &mask)
62
63     ...
64
65     //TODO0 create mask for values for an obtained tmp values, which are large
66     than 0.5 and
67     //TODO0 put all values smaller than 0.5 from tmp to the array "output" at the
68     places given by indices "index"
69     // Use void scatter (float *array, const uint_v &indexes, const float_m
70     &mask) const

```

	Typical output
	Indices: 25 96 76 42 Gather without masking: results are WRONG. -4.67253e+33 is not equal to 0.988475 4.59163e-41 is not equal to 0.912037 0 is not equal to 0.80679 Gather with masking: results are WRONG. Scatter with masking: results are correct.

3_Vc/3_RandomAccess: solution

(1) To gather data from the given places in an array, one needs to create a **uint_v** vector with corresponding indexes, then the **gather** function can be applied directly.

	Part of the source code of RandomAccess_solution.cpp
53	<code>uint_v ind(index);</code>
54	<code>tmp.gather(input, ind);</code>

(2) To gather data under the given condition, one needs to create a mask with a corresponding type. Since the data is **float**, the **float_m** mask must be created, then the **gather** function with the mask parameter is applied directly. To ensure that the entries, which were masked out, have meaningful values, we need to initialise the output variable before gathering.

	Part of the source code of RandomAccess_solution.cpp
67	<code>float_m mask = tmp > 0.5f;</code>
68	<code>float_v tmp2(Vc::Zero);</code>
69	<code>tmp2.gather(input, ind, mask);</code>

(3) To scatter data under the given condition one needs to create a mask with a corresponding type. Since the data is **float**, the same mask can be used as in (2), then the **scatter** function with the mask parameter is applied directly.

	Part of the source code of RandomAccess_solution.cpp
93	<code>mask = tmp < 0.5f;</code>
94	<code>tmp.scatter(output, ind, mask);</code>
	Typical output after solution
	Indices: 98 40 24 14 Gather without masking: results are correct. Gather with masking: results are correct. Scatter with masking: results are correct.