

High Performance Computer Architectures Practical Course - Exercise 3 -

Tutorium 1

David Jordan (6260776)
Florian Rüffer (7454628)
Michael Sanjatin (7485765)

May 5, 2023

1 Neural Networks and SIMD

1

2

The task for this exercise is to add SIMD (Single Instruction, Multiple Data) support to our neural network functions. We will start of with the simd-ized version of the ReLU activation function.

The first noticable detail is the changed input. We input type is "fvec". Fvec is a SIMD-packed set of values, to be exact 4 packed 32-bit floating point numbers. Once again we create a result vector, the same size as the input. Subsequently we loop for 1.) every element of fvec (here this is 4) and 2.) the number of elements of the input. As a result, we loop over every element of the input.

Finally, we check the conditions associated with the ReLU activation function for every of those elements. The defintion of ReLU is provided in the Appendix (5.1)

```
1      std::vector<fvec> MLPMath::applyReLU(std::vector<
      fvec>& input) {
2
3      std::vector<fvec> result(input.size());
4
5      for (std::size_t iv = 0; iv < fvecLen; iv++) {
6          for (std::size_t i = 0; i < input.size(); i++)
7              {
8                  if (input[i][iv] > 0.0f) {
9                      result[i][iv] = input[i][iv];
10                 }
11                 else {
12                     result[i][iv] = 0.0f;
13                 }
14             }
15
16      return result;
17 }
```

File 1: SIMD-ized ReLU

Another essential part is the `backPropActivation` function. Here we provide just the code snippet, which specifically implements ReLU.

In this case there are not many significant changes. We want to ensure, that this function will execute properly if provided with a `fvec` vector input. For this to work, line 4 is the most essential one to understand. Here the `fvec`-type overloads the `>`-operator. To understand the underlying functionality, please view the following example.

"rawOutput" contains the elements: [1.0, -1.0, 2.0, -0.5]

"zero" contains the elements [0.0, 0.0, 0.0, 0.0]

Overloading will produce a result like this: [1.0 > 0.0, -1.0 > 0.0, 2.0 > 0.0, -0.5 > 0]

This checks the condition for every element and can be abbreviated into the following form: [0xFFFFFFFF, 0x00000000, 0xFFFFFFFF, 0x00000000] (1 - True & 0 - False)

```
1      case 1:{
2          fvec zero = fvec(0.0f);
3          for (std::size_t i = 0; i < rDelta.size(); i
4              ++){
5              rDelta[i] = (rawOutput[i] > zero) * rDelta
6                  [i];
7          }
8          break;
9      }
```

File 2: `backPropActivation`

Finally we need to incorporate our SIMD-ized ReLU function into the apply-Activation function (Line 10-14).

```

1      std::vector<fvec> MLPNet::applyActivation(std::
      vector<fvec>& input) {
2
3      std::vector<fvec> output(input.size());
4      switch (activationType_) {
5          case 0: {// TanH
6              output = MLPMath::applyTanH(input);
7              return output;
8              break;
9          }
10         case 1: {// SIMD-ized ReLU
11             output = MLPMath::applyReLU(input);
12             return output;
13             break;
14         }
15         default: {
16             return output;
17             break;
18         }
19     }
20 }
```

File 3: applyActivation

3

2 Matrix

In this task we need to speed up Matrix calculations using SIMD. As displayed below, we want to go from matrix a to matrix c.

$$a = \begin{bmatrix} a_0 & a_1 & a_2 \\ a_3 & a_4 & a_5 \\ a_6 & a_7 & a_8 \end{bmatrix}$$

$$c = \begin{bmatrix} \sqrt{a_0} & \sqrt{a_1} & \sqrt{a_2} \\ \sqrt{a_3} & \sqrt{a_4} & \sqrt{a_5} \\ \sqrt{a_6} & \sqrt{a_7} & \sqrt{a_8} \end{bmatrix}$$

First of all we need to change the input & output matrices (output for scalar computation stays the same). We do this by changing the memory alignment of the values of the matrices. This is a requirement for SIMD-ized operations.

```

1 float a[N][N] __attribute__((aligned(16))); //
   input array
2 float c[N][N]; // output array for scalar computations
3 float c_simd[N][N] __attribute__((aligned(16))); //
   output array for SIMD computations

```

File 4: Matrix.cpp

The loops for the computation part remain mostly unchanged to the scalar version. The inner loop runs N -times but our loop variable increases by 4, because each vector has 4 elements and those can be calculated in parallel. We loop for $N_{iter} - 1$ times, to neglect memory reading time. Subsequently we loop over matrix a to transform it into matrix c ($N \times N$ for rows and columns). The distinct section ranges from line 5 to line 7. We define two new vectors, called $aVec$ and $cVec$, which suppose to be the vector counter parts to the scalar versions of matrix a and matrix c . At the beginning those were initialized with float values, so we use the `reinterpret_cast` command to allow treating a variables memory representation as another type. In this case we substitute float for `fvec`. Finally we just plug those values into the template function to calculate the root (Line 7).

```

1 TStopwatch timerSIMD;
2 for( int ii = 0; ii < Niter; ii++ )
3     for( int i = 0; i < N; i++ ) {
4         for( int j = 0; j < N; j++ ) {
5             fvec &aVec = reinterpret_cast<fvec>(a[i][
                        j]);
6             fvec &cVec = reinterpret_cast<fvec>(
                        c_simd[i][j]);
7             cVec = f(aVec);
8         }
9     }
10 timerSIMD.Stop();

```

File 5: Matrix.cpp

We can compute 4 values in parallel and therefore can expect a speed-up factor of 4. Due to running environment deviation, the speed-up factor varies.

```

Time scalar: 355.03 ms
Time SIMD: 405.338 ms, speed up 0.875886
SIMD and scalar results are the same.

```


Figure 1: Output

3 Quadratic Equation

4 CheckSum

5 Appendix

5.1 ReLU



$$\text{ReLU}(x) = \begin{cases} x, & \text{if } x \geq 0 \\ 0, & \text{otherwise} \end{cases}$$



Figure 2: Add caption